

DEVA ANAND

devaanand@umass.edu • (413) 315-1715 • devaanand.com • linkedin.com/in/devaaa
• github.com/Deva-1903

June 30, 2026

Hiring Team, AI Center of Excellence
Black Box
Plano, Texas

Dear Hiring Team,

I am applying for the AI Engineer Intern role in the AI Center of Excellence. I am a full-time MS Computer Science student at UMass Amherst (graduating May 2027), and before grad school I spent two years as a backend engineer shipping production systems. Your description of building enterprise-grade, production-ready GenAI that plugs into real business platforms is exactly the kind of applied AI work I want to do, so I wanted to write directly.

Most of my background is building the things this role asks for. At Wysa I engineered a multi-tenant NHS eTriage backend in Node.js and MongoDB for 8+ UK clinical clients, processing 10,000+ monthly submissions through REST integrations with Mayden's iaptus enterprise platform, so integrating AI into systems that already carry sensitive data and external dependencies is familiar territory. I also architected a full-stack React and Node.js data-annotation platform used daily by our AI team that cut manual labeling by 100+ hours per month, and I owned the security side of the stack, including AWS KMS encryption and remediating 20+ VAPT findings. Earlier, at Cario, I shipped a production LLM feature end to end using a 4-bit quantized Llama-7B model via llama.cpp with per-persona prompts, async queued generation, and content-safety filtering.

On the Generative AI side, I have built the RAG and agent patterns central to this team. My AgenticSearch project is a multi-stage agentic service that turns open-ended queries into structured entity tables through typed retrieval planning, reranking, and deterministic extraction with LLM fallback, returning cell-level provenance and hardened with 150+ tests. In a KG-guided RAG pipeline for multi-hop QA, I built the multi-view retrieval component (query decomposition, RRF fusion, cross-encoder reranking, MMR) and lifted supporting-fact recall by 2.68% over the baseline. I work in Python, TypeScript, FastAPI, Node.js, and React daily, and I use LLM APIs and AI coding tools as accelerators while keeping behavior grounded in tests and provenance, an approach I first practiced in a measured GitHub Copilot pilot at Wysa.

What draws me to this role is the emphasis on doing applied AI properly: production-ready, secure, and integrated into the business rather than left as a demo. I like work where a small mistake actually costs something, and I learn fastest when I am a little out of my depth, so a 6 to 12 month internship building real enterprise AI use cases is exactly where I want to be. I would welcome the chance to talk about how I can contribute to the CoE.

Thank you for your time and consideration.

Sincerely,
Deva Anand