

DEVA ANAND

devaanand@umass.edu • (413) 315-1715 • devaanand.com • linkedin.com/in/devaaa
• github.com/Deva-1903

July 2, 2026

Hiring Team, AI
Jaan Health / Phamily
New York, NY

Dear Hiring Team,

I am applying for the Machine Learning Intern (R&D) role. What drew me in is that Jaan Health is doing applied AI where it is hardest and matters most: turning chronic-care management from reactive to proactive, on real clinical data, with model quality tied directly to patient outcomes. I spent two years building production healthcare software before grad school, and I now work on LLM, retrieval, and evaluation systems in my MS, so a role at the intersection of applied research and production healthcare systems is exactly where I want to be.

At Wysa, I engineered a multi-tenant NHS clinical eTriage backend serving 8+ UK clinical clients and 10,000+ monthly submissions, and I handled sensitive clinical data in a regulated setting, AWS KMS encryption for PII, retention policies, and security hardening. I also built a full-stack annotation platform used daily by our AI team to prepare model training data, which cut manual annotation by 100+ hours per month. So the parts of this role that involve messy real-world healthcare data, regulated constraints, and getting models the data they need are familiar territory, not new ground. My first publication was also in this space: a deep transfer-learning approach to Alzheimer's classification at IEEE CONIT 2023.

On the ML side, my recent work is squarely in NLP, LLMs, and retrieval. I built AgenticSearch, a multi-stage agentic RAG service with deterministic extraction, cell-level provenance, and 150+ tests, deployed as a live service. I led the multi-view retrieval component of a KG-guided RAG pipeline that lifted supporting-fact recall by 2.68% over its baseline. And with AutoEval I built an evaluation-improvement agent with anti-Goodharting guardrails, because I care a lot about evaluation that does not quietly fool itself, which feels essential when model performance is connected to patient outcomes rather than a leaderboard. Earlier, at Cario, I shipped a production conversational LLM feature end to end, so I have taken model work from prototype to something people actually used.

Your principles of Care, Curiosity, and Craftsmanship resonate with how I already work. I like problems where a small mistake actually costs something, and I am most motivated when research has to become a real system that holds up in production. I would be glad to bring that to the AI team at Jaan Health.

Thank you for your time and consideration.

Sincerely,
Deva Anand