

Deva Anand

(413) 315-1715 | devaanand@umass.edu | devaanand.com | linkedin.com/in/devaaa
| github.com/Deva-1903

May 24, 2026

Snowflake University Recruiting

Re: Software Engineer Intern (Infrastructure Automation) — Fall 2026

Dear Snowflake Hiring Team,

I am a Master's student in Computer Science at UMass Amherst (graduating May 2027), applying for your Fall 2026 Software Engineer Internship in Infrastructure Automation. What draws me to this role is its leverage: building the automation frameworks, tooling, and release infrastructure that let everyone else ship faster and more reliably. I would rather make a whole engineering org more productive than ship one more feature, and a data platform that real customers depend on is exactly where that work matters most.

That instinct shows up in what I have already built. At Wysa I built an internal annotation platform that automated a manual, spreadsheet-based workflow into secure role-based APIs, cutting 100+ hours of repetitive operations a month and lifting throughput 60% for the team that relied on it. I also owned reliability for a multi-tenant NHS backend serving 8+ UK clients, integration tests, authentication hardening, 20+ VAPT remediations, and release monitoring, because trustworthy software is the whole point on sensitive data. More recently I optimized a Spark ETL pipeline over a 1.4-billion-row dataset: profiling stage-level timings in the Spark UI, enabling adaptive query execution and skew-join handling, cutting runtime by roughly 25%, and turning the fixes into a repeatable optimization playbook rather than a one-off. That mix of large-scale data, performance, and repeatable automation maps closely to your Engineering Systems and Quality & Release Engineering work.

Your description of treating AI as a high-trust collaborator matches how I work day to day. I lean on tools like Claude Code and Cursor as accelerators, but I keep system behavior grounded in tests, CI, and verifiable evidence. My agentic web-discovery project, for instance, pairs LLM extraction with deterministic validation, cell-level provenance, and 150+ automated tests. I have not used Jenkins, Ansible, or Terraform in production yet, but I learn fastest when I am a little out of my depth, and picking up new infrastructure tooling quickly is exactly the kind of challenge I am looking for. I would be glad to bring that to a real system Snowflake customers rely on.

Thank you for your time and consideration.

Sincerely,

Deva Anand